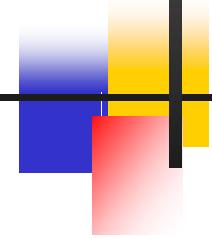
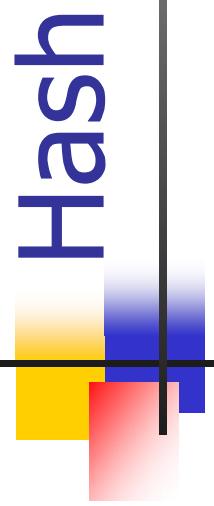


Hashing & Hash Tables

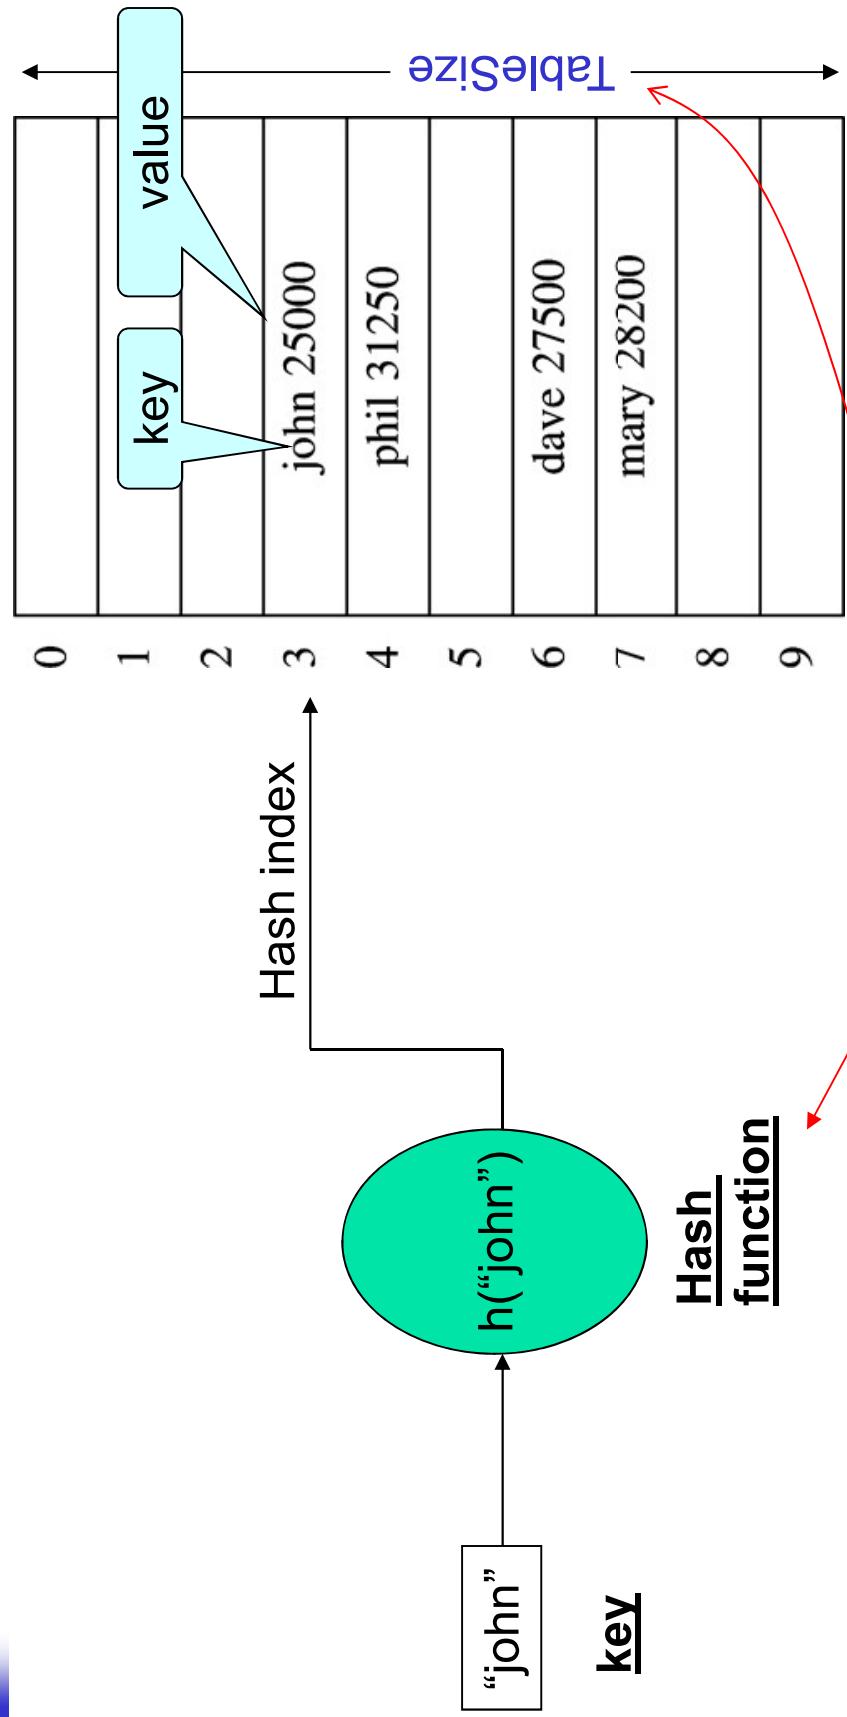


Overview

- Hash Table Data Structure : Purpose
 - *To support insertion, deletion and search in average-case constant time*
 - Assumption: Order of elements irrelevant
 - => data structure *not* useful for if you want to maintain and retrieve some kind of an order of the elements
- Hash function
 - *Hash["string key"] ==> integer value*
- Hash table ADT
 - Implementations, Analysis, Applications



Hash table: Main components



Hash table
(implemented as a vector)

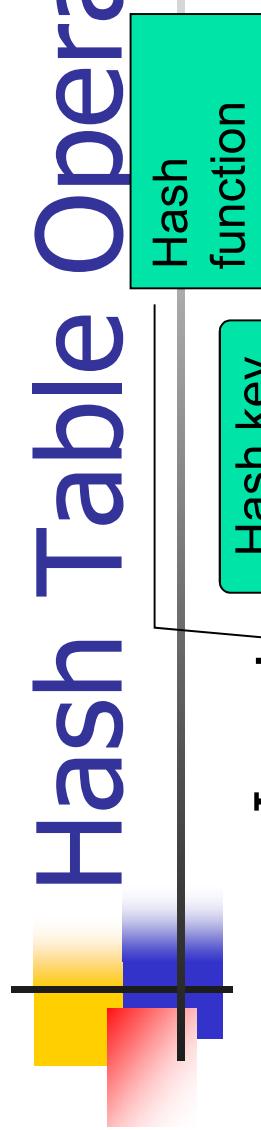
Hash Table

- Hash table is an array of fixed size TableSize
- Array elements indexed by a **key**, which is mapped to an array index (0...TableSize-1)
- Mapping (hash function) h from key to index
 - E.g., $h(\text{"john"}) = 3$

0	
1	
2	
3	john 25000
4	phil 31250
5	
6	dave 27500
7	mary 28200
8	
9	

key Element value

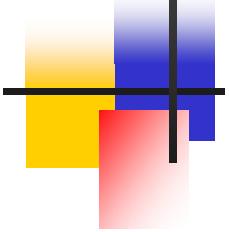
Hash Table Operations



- Insert
 - $T[h(\text{"john"})] = <\text{"john"}, 25000>$
 - Delete
 - $T[h(\text{"john"})] = \text{NULL}$
 - Search
 - $T[h(\text{"john"})]$ returns the element hashed for "john"
- What happens if $h(\text{"john"}) == h(\text{"joe"})$?
"collision"**

0	
1	
2	
3	john 25000
4	phil 31250
5	
6	dave 27500
7	mary 28200
8	
9	

Factors affecting Hash Table Design



- Hash function
- Table size
 - Usually fixed at the start
- Collision handling scheme

Hash Function

- A hash function is one which maps an element's key into a valid hash table index
 - $h(\text{key}) \Rightarrow \text{hash table index}$

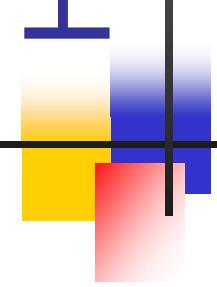
Note that this is (slightly) different from saying:

$h(\text{string}) \Rightarrow \text{int}$

- Because the key can be of any type
 - E.g., " $h(\text{int}) \Rightarrow \text{int}$ " is also a hash function!
 - But also note that any type can be converted into an equivalent string form

$h(\text{key}) \Rightarrow \text{hash table index}$

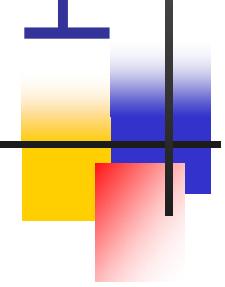
Hash Function Properties



- A hash function maps key to integer
- **Constraint:** Integer should be between [0, TableSize-1]
- A hash function can result in a many-to-one mapping (causing collision)
- Collision occurs when hash function maps two or more keys to same array index
- Collisions cannot be avoided but its chances can be reduced using a "good" hash function

$h(\text{key}) \Rightarrow \text{hash table index}$

Hash Function Properties



- A “good” hash function should have the properties:

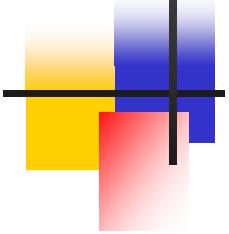
1. Reduced chance of collision

Different keys should ideally map to different indices

Distribute keys uniformly over table

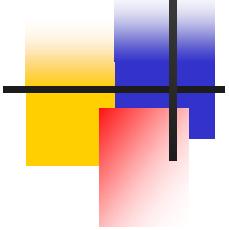
2. Should be fast to compute

Hash Function - Effective use of table size



- Simple hash function (assume integer keys)
 - $h(\text{Key}) = \text{Key} \bmod \text{TableSize}$
- For random keys, $h()$ distributes keys evenly over table
 - What if TableSize = 100 and keys are ALL multiples of 10?
 - Better if TableSize is a prime number

Different Ways to Design a Hash Function for String Keys



A very simple function to map strings to integers:

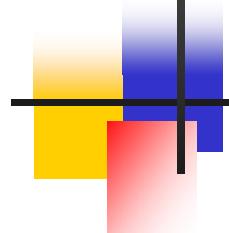
- Add up character ASCII values (0-255) to produce integer keys

- E.g., "abcd" = 97+98+99+100 = 394
- ==> $h("abcd") = 394 \% \text{TableSize}$

Potential problems:

- Anagrams will map to the same index
 - $h("abcd") == h("dbac")$
- Small strings may not use all of table
 - $\text{Strlen}(S) * 255 < \text{TableSize}$
- Time proportional to length of the string

Different Ways to Design a Hash Function for String Keys



Approach 2

- Treat first 3 characters of string as base-27 integer (26 letters plus space)
 - $\text{Key} = S[0] + (27 * S[1]) + (27^2 * S[2])$
- Better than approach 1 because ... ?

Potential problems:

- Assumes first 3 characters randomly distributed
 - Not true of English
-
- collision
- Apple
Apply
Appointment
Apricot

Different Ways to Design a Hash Function for String Keys

Approach 3

Use all N characters of string as an N-digit base-K number

- Choose K to be prime number larger than number of different digits (characters)
 - I.e., K = 29, 31, 37
- If L = length of string S, then

$$h(S) = \left[\sum_{i=0}^{L-1} S[L-i-1] * 37^i \right] \bmod TableSize$$

- Use Horner's rule to compute h(S)
- Limit L for long strings

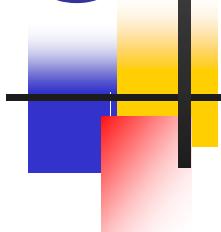
```
1  /**
2   * A hash routine for string objects.
3   */
4  int hash( const string & key, int tableSize )
5  {
6      int hashVal = 0;
7
8      for( int i = 0; i < key.length( ); i++ )
9          hashVal = 37 * hashVal + key[ i ];
10
11     hashVal %= tableSize;
12     if( hashVal < 0 )
13         hashVal += tableSize;
14
15     return hashVal;
16 }
```

Problems:

- potential overflow
- larger runtime

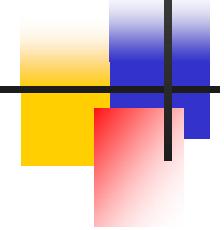
“Collision resolution techniques”

Techniques to Deal with Collisions



- Chaining
- Open addressing
- Double hashing
- Etc.

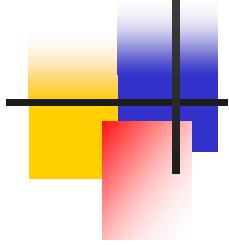
Resolving Collisions



- What happens when $h(k_1) = h(k_2)$?
 - ==> collision !
- Collision resolution strategies
 - ***Chaining***
 - Store colliding keys in a linked list at the same hash table index
 - ***Open addressing***
 - Store colliding keys elsewhere in the table

Collision resolution technique #1

Chaining



Chaining strategy: maintains a linked list at every hash index for collided elements

Insertion sequence: { 0 1 4 9 16 25 36 49 64 81 }

- Hash table T is a vector of linked lists
 - Insert element at the head (as shown here) or at the tail
 - Key k is stored in list at $T[h(k)]$
 - E.g., TableSize = 10
 - $h(k) = k \bmod 10$
 - Insert first 10 perfect squares
-
- | | | | | | | | | | | |
|---|----|--|--|----|----|----|---|----|----|--|
| 0 | 81 | | | 64 | 25 | 36 | 4 | 16 | 49 | |
| 1 | | | | | | | | | | |
| 2 | | | | | | | | | | |
| 3 | | | | | | | | | | |
| 4 | | | | | | | | | | |
| 5 | | | | | | | | | | |
| 6 | | | | | | | | | | |
| 7 | | | | | | | | | | |
| 8 | | | | | | | | | | |
| 9 | | | | | | | | | | |

Implementation of Chaining Hash Table

```
1 template <typename HashedObj>
2 class HashTable
3 {
4     public:
5         explicit HashTable( int size = 101 );
6
7     bool contains( const HashedObj & x ) const;
8
9     void makeEmpty( );
10    void insert( const HashedObj & x );
11    void remove( const HashedObj & x );
12
13 private:
14     vector<list<HashedObj>> theLists; // The array of Lists
15     int currentSize;
16
17     void rehash( );
18     int myhash( const HashedObj & x ) const;
19 };
20
21 int hash( const string & key );
22 int hash( int key );
```

Vector of linked lists
(this is the main hashtable)

Current #elements in
the hashtable

Hash functions for
integers and string
keys

Implementation of Chaining Hash Table

```
1  int myhash( const HashedObj & x ) const
2  {
3      int hashVal = hash( x );
4
5      hashVal %= theLists.size( );
6
7      if( hashVal < 0 )
8          hashVal += theLists.size( );
9
10 }
```

This is the hashtable's current capacity (aka. "table size")

This is the hash table index for the element x

```

1   bool insert( const HashedObj & x )
2   {
3       list<HashedObj> & whichList = theLists[ myhash( x ) ];
4       if( find( whichList.begin( ), whichList.end( ), x ) != whichList.end( ) )
5           return false;
6       whichList.push_back( x );
7       // Rehash; see Section 5.5
8       if( ++currentSize > theLists.size( ) )
9           rehash( );
10      currentSize = theLists.size();
11      return true;
12  }
13

```

Duplicate check

Later, but essentially
resizes the hashtable if its
getting crowded

```

1 void makeEmpty( )
2 {
3     for( int i = 0; i < theLists.size( ); i++ )
4         theLists[ i ].clear( );
5     }
6
7     bool contains( const HashedObj & x ) const
8     {
9         const list<HashedObj> & whichList = theLists[ myhash( x ) ];
10        return find( whichList.begin( ), whichList.end( ), x ) != whichList.end( );
11    }
12
13    bool remove( const HashedObj & x )
14    {
15        list<HashedObj> & whichList = theLists[ myhash( x ) ];
16        list<HashedObj>::iterator itr = find( whichList.begin( ), whichList.end( ), x );
17
18        if( itr == whichList.end( ) )
19            return false;
20
21        whichList.erase( itr );
22        --currentSize;
23        return true;
24    }

```

Each of these operations takes time linear in the length of the list at the hashed index location

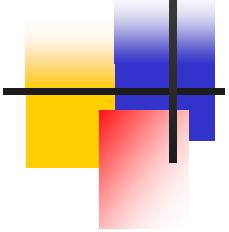
```

1 // Example of an Employee class
2 class Employee
3 {
4     public:
5         const string & getName( ) const
6             { return name; }
7
8     bool operator==( const Employee & rhs ) const
9         { return getName( ) == rhs.getName( ); }
10    bool operator!=( const Employee & rhs ) const
11        { return !( *this == rhs); }
12
13    // Additional public members not shown
14
15    private:
16        string name;
17        double salary;
18        int seniority;
19
20    // Additional private members not shown
21 };
22
23 int hash( const Employee & item )
24 {
25     return hash( item.getName( ) );
26 }
```

All hash objects must define == and != operators.

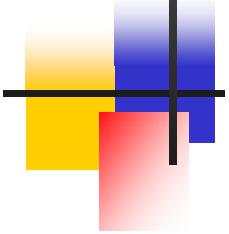
Hash function to handle Employee object type

Collision Resolution by Chaining: Analysis



- **Load factor** λ of a hash table T is defined as follows:
 - N = number of elements in T ("current size")
 - M = size of T ("table size")
 - $\lambda = N/M$ ("load factor")
 - i.e., λ is the average length of a chain
- **Unsuccessful search time:** $O(\lambda)$
 - Same for insert time
- **Successful search time:** $O(\lambda/2)$
- **Ideally,** want $\lambda \leq 1$ (not a function of N)

Potential disadvantages of Chaining



Linked lists could get long

- Especially when N approaches M
- Longer linked lists could negatively impact performance

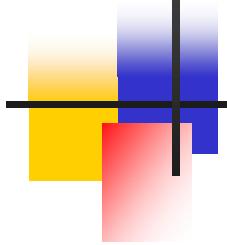
More memory because of pointers

Absolute worst-case (even if $N << M$):

- All N elements in one linked list!
- Typically the result of a bad hash function

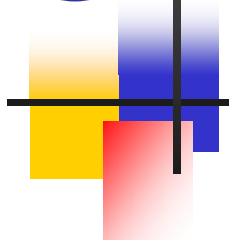
Collision resolution technique #2

Open Addressing



Collision Resolution by Open Addressing

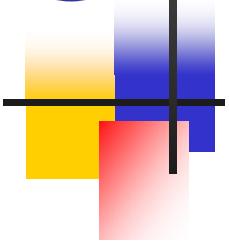
An “inplace” approach



When a collision occurs, look elsewhere in the table for an empty slot

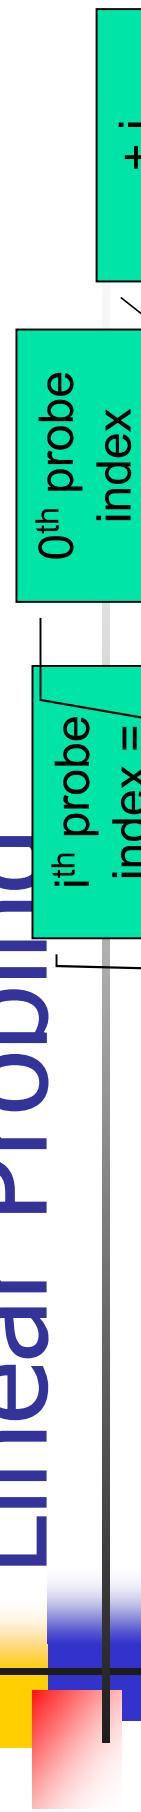
- Advantages over chaining
 - No need for list structures
 - No need to allocate/deallocate memory during insertion/deletion (slow)
- Disadvantages
 - Slower insertion – May need several attempts to find an empty slot
 - Table needs to be bigger (than chaining-based table) to achieve average-case constant-time performance
 - Load factor $\lambda \approx 0.5$

Collision Resolution by Open Addressing



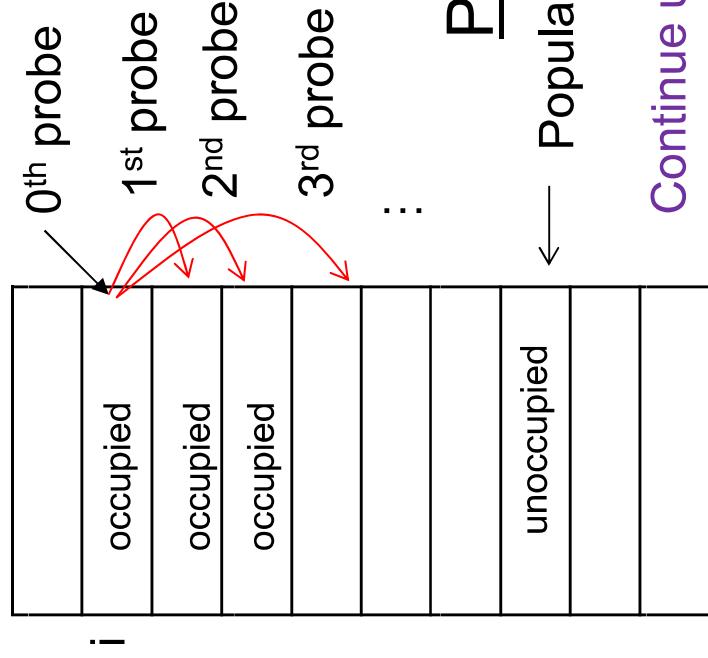
- A "probe sequence" is a sequence of slots in hash table while searching for an element x
 - $h_0(x), h_1(x), h_2(x), \dots$
 - Needs to visit each slot exactly once
 - Needs to be repeatable (so we can find/delete what we've inserted)
- Hash function
 - **$h_i(x) = (h(x) + f(i)) \bmod \text{TableSize}$**
 - $f(0) = 0$
 - $f(i)$ is "*the distance to be traveled relative to the 0th probe position, during the i^{th} probe*".

Linear Probing

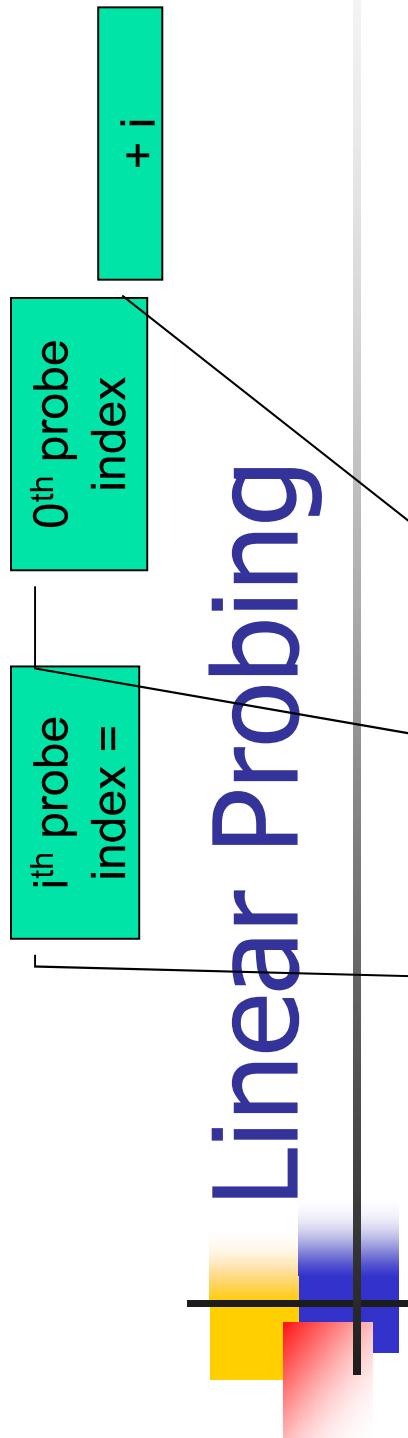


Linear probing:

- $f(i)$ = is a linear function of i ,



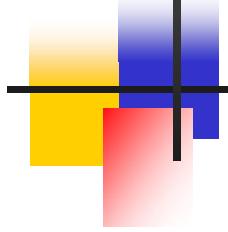
Continue until an empty slot is found
#failed probes is a measure of performance



- $f(i) \neq i$ is a linear function of i , e.g., $f(i) = i$
- $h_i(x) = (h(x) + i) \text{ mod } \text{TableSize}$

- Probe sequence: +0, +1, +2, +3, +4, ...
- Example: $h(x) = x \bmod \text{TableSize}$
 - $h_0(89) = (h(89) + f(0)) \bmod 10 = 9$
 - $h_0(18) = (h(18) + f(0)) \bmod 10 = 8$
 - $h_0(49) = (h(49) + f(0)) \bmod 10 = 9$ (X)
 - $h_1(49) = (h(49) + f(1)) \bmod 10$
 - $= (h(49) + 1) \bmod 10 = 0$

Linear Probing Example



Insert sequence: 89, 18, 49, 58, 69

	Empty Table	After 89	After 18	After 49	After 58	After 69	
0		49		49		49	
1			58		58		
2				69			
3							
4							
5							
6							
7							
8			18		18		
9		89		89		89	

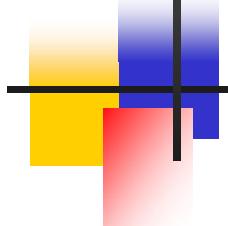
#unsuccessful
probes:

7

Cpt S 223. School of EECS, WSU

total 30

Linear Probing: Issues

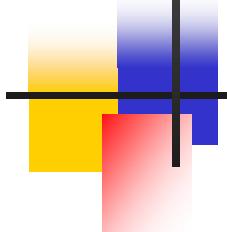


Probe sequences can get longer with time

Primary clustering

- Keys tend to cluster in one part of table
- Keys that hash into cluster will be added to the end of the cluster (making it even bigger)
- Side effect: Other keys could also get affected if mapping to a crowded neighborhood

Linear Probing: Analysis

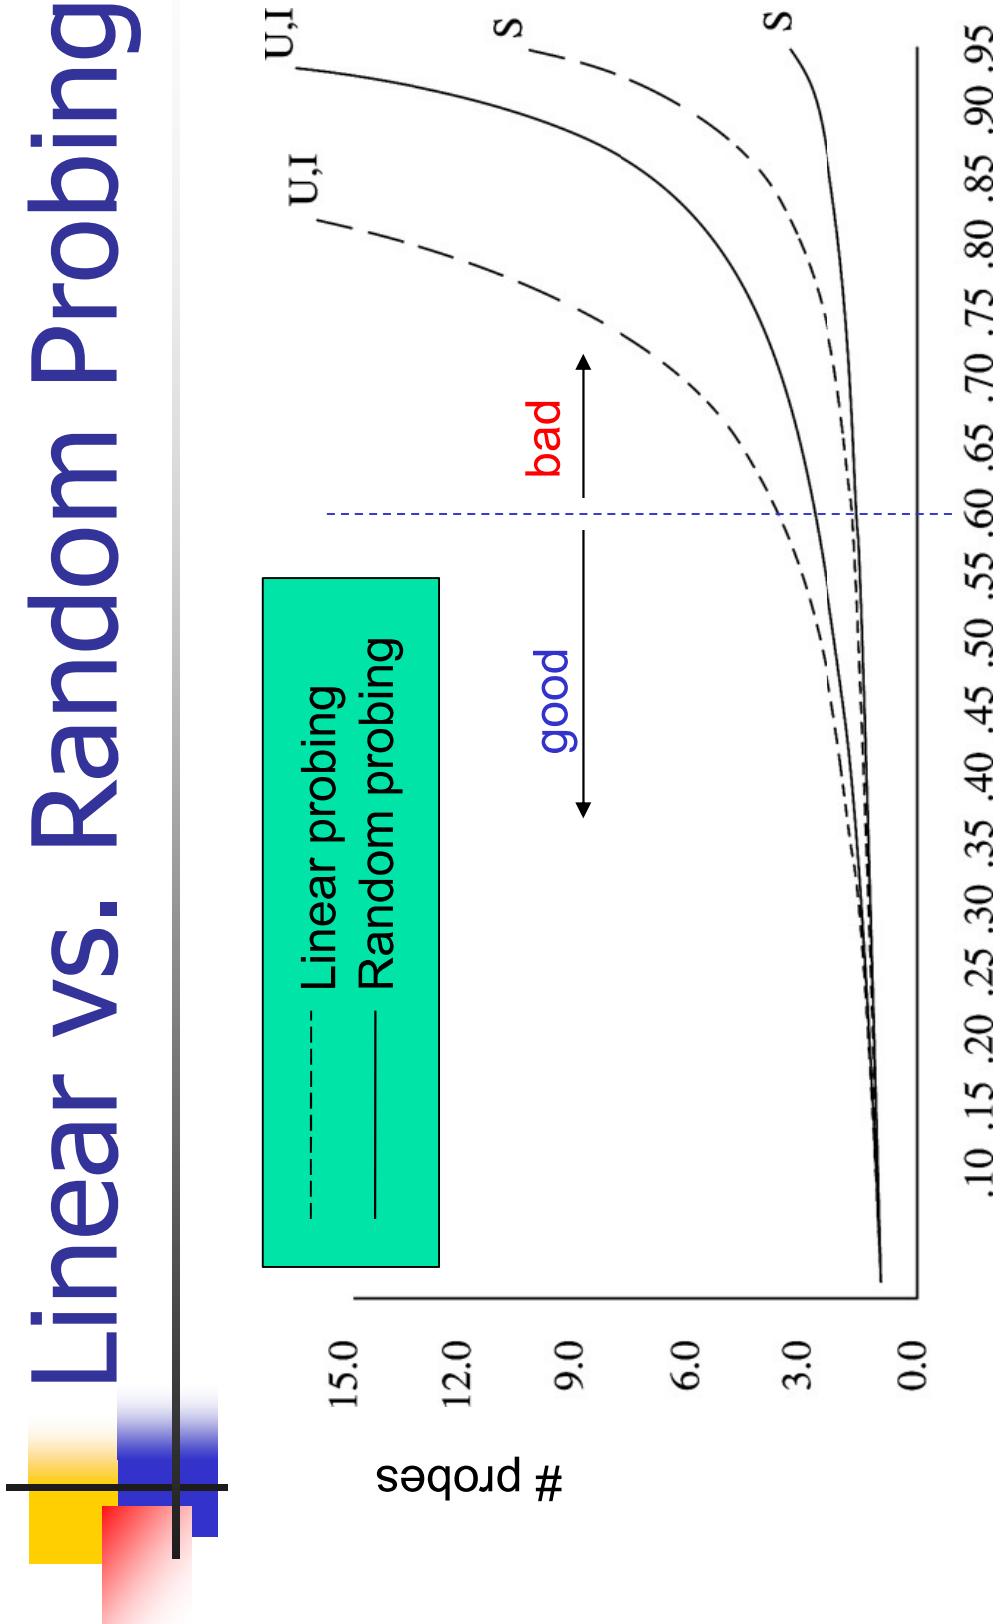


- Expected number of probes for insertion or unsuccessful search
$$\frac{1}{2} \left(1 + \frac{1}{(1-\lambda)^2} \right)$$
- Example ($\lambda = 0.5$)
 - Insert / unsuccessful search
 - 2.5 probes
 - Successful search
 - 1.5 probes
- Expected number of probes for successful search
$$\frac{1}{2} \left(1 + \frac{1}{(1-\lambda)} \right)$$
- Example ($\lambda = 0.9$)
 - Insert / unsuccessful search
 - 50.5 probes
 - Successful search
 - 5.5 probes

Random Probing: Analysis

- Random probing does not suffer from clustering
- Expected number of probes for insertion or unsuccessful search:
$$\frac{1}{\lambda} \ln \frac{1}{1-\lambda}$$
- Example
 - $\lambda = 0.5$: 1.4 probes
 - $\lambda = 0.9$: 2.6 probes

Linear VS. Random Probing

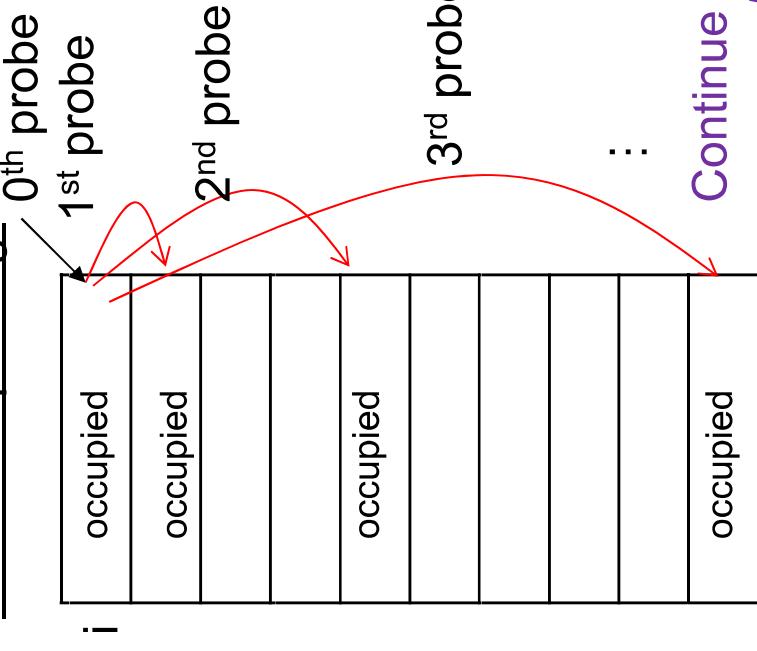


U - unsuccessful search
S - successful search
I - insert

Load factor λ

Quadratic Probing

Quadratic probing:



- Avoids primary clustering

- $f(i)$ is quadratic in i

e.g., $f(i) = i^2$

$$h_i(x) = (h(x) + i^2) \bmod \text{TableSize}$$

- Probe sequence:
- +0, +1, +4, +9, +16, ...

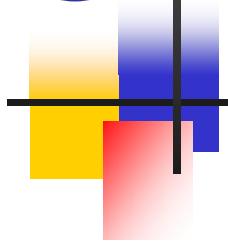
Continue until an empty slot is found
#failed probes is a measure of performance

Quadratic Probing

- Avoids primary clustering
- $f(i)$ is quadratic in I , e.g., $f(i) = i^2$
 - $h_i(x) = (h(x) + i^2) \bmod \text{TableSize}$
 - Probe sequence: +0, +1, +4, +9, +16, ...
- Example:
 - $h_0(58) = (h(58)+f(0)) \bmod 10 = 8 \text{ (X)}$
 - $h_1(58) = (h(58)+f(1)) \bmod 10 = 9 \text{ (X)}$
 - $h_2(58) = (h(58)+f(2)) \bmod 10 = 2$

Q) Delete(49), Find(69) - is there a problem?

Quadratic Probing Example



Insert sequence: 89, 18, 49, 58, 69

	Empty Table	After 89	After 18	After 49	After 58	After 69	
0				49	+1 ² 49	49	+1 ²
1					58	+2 ² 58	
2						69	+2 ²
3							
4							
5							
6							
7							+0 ²
8							
9							

#unsuccessful
probes:

0 0 1 2 2 5

Quadratic Probing: Analysis

- Difficult to analyze
- Theorem 5.1
 - New element can always be inserted into a table that is at least half empty and TableSize is prime
 - Otherwise, may never find an empty slot, even if one exists
- Ensure table never gets half full
 - If close, then expand it

Quadratic Probing

- May cause “secondary clustering”

- Deletion

- Emptying slots can break probe sequence and could cause find stop prematurely
- Lazy deletion
 - Differentiate between empty and deleted slot
 - When finding skip and continue beyond deleted slots
 - If you hit a non-deleted empty slot, then stop find procedure returning “not found”
- May need compaction at some time

Quadratic Probing: Implementation

```
1 template <typename HashedObj>
2 class HashTable
3 {
4     public:
5         explicit HashTable( int size = 101 );
6
7     bool contains( const HashedObj & x ) const;
8
9     void makeEmpty( );
10    bool insert( const HashedObj & x );
11    bool remove( const HashedObj & x );
12
```

Quadratic Probing: Implementation

```
13 enum EntryType { ACTIVE, EMPTY, DELETED };
14
15 private:
16     struct HashEntry {
17         HashedObj element;
18         EntryType info;
19     };
20
21     HashEntry( const HashedObj & e = HashedObj( ), EntryType i = EMPTY )
22         : element( e ), info( i ) { }
23     };
24
25     vector<HashEntry> array;
26     int currentSize;
27
28     bool isActive( int currentPos ) const;
29     int findPos( const HashedObj & x ) const;
30     void rehash( );
31     int myhash( const HashedObj & x ) const;
32 }
```

Lazy deletion

Quadratic Probing: Implementation

```
1 explicit HashTable( int size = 101 ) : array( nextPrime( size ) )
2 { makeEmpty( ); }
3
4 void makeEmpty( )
5 {
6     currentSize = 0;
7     for( int i = 0; i < array.size( ); i++ )
8         array[ i ].info = EMPTY;
9 }
```

Ensure table
size is prime

Quadratic Probing: Implementation

```
1  bool contains( const HashedObj & x ) const
2  { return isActive( findPos( x ) ); }

3
4  int findPos( const HashedObj & x ) const
5  {
6      int offset = 1;
7      int currentPos = myhash( x );
8
9      while( array[ currentPos ].info != EMPTY &&
10         array[ currentPos ].element != x )
11      {
12          currentPos += offset; // Compute ith probe
13          offset += 2;
14          if( currentPos >= array.size( ) )
15              currentPos -= array.size( );
16      }
17
18      return currentPos;
19  }

20
21  bool isActive( int currentPos ) const
22  { return array[ currentPos ].info == ACTIVE; }
```

Find

Skip DELETED;
No duplicates

Quadratic probe
sequence (really)

Quadratic Probing:

Implementation

```
1  bool insert( const HashedObj & x )
2  {
3      // Insert x as active
4      int currentPos = findPos( x );
5      if( isActive( currentPos ) )
6          return false;
7
8      array[ currentPos ] = HashEntry( x, ACTIVE );
9
10     // Rehash; see Section 5.5
11     if( ++currentSize > array.size( ) / 2 )
12         rehash( );
13
14     return true;
15 }
16
17 bool remove( const HashedObj & x )
18 {
19     int currentPos = findPos( x );
20     if( !isActive( currentPos ) )
21         return false;
22
23     array[ currentPos ].info = DELETED;
24     return true;
25 }
```

The diagram illustrates the flow of control in the quadratic probing implementation. Three green callout boxes point to specific parts of the code:

- An arrow points from line 3 to a box labeled "Insert".
- An arrow points from line 5 to a box labeled "No duplicates".
- An arrow points from line 11 to a box labeled "Remove".
- An arrow points from line 23 to a box labeled "No deallocation needed".

Double Hashing: keep two hash functions h_1 and h_2

- Use a second hash function for all tries 1 other than 0:
 $f(i) = i * h_2(x)$
- Good choices for $h_2(x)$?
 - Should never evaluate to 0
 - $h_2(x) = R - (x \bmod R)$
 - R is prime number less than TableSize
- Previous example with $R=7$
 - $h_0(49) = (h(49)+f(0)) \bmod 10 = 9 (X)$
 - $h_1(49) = (h(49)+1*(7 - 49 \bmod 7)) \bmod 10 = 6$

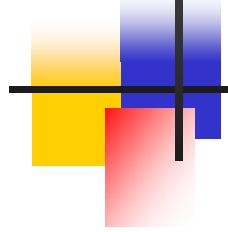
Double Hashing Example

	Empty Table	After 89	After 18	After 49	After 58	After 69
0						69
1						
2						
3					58	58
4						
5						
6				49	49	49
7						
8			18	18	18	18
9		89	89	89	89	89

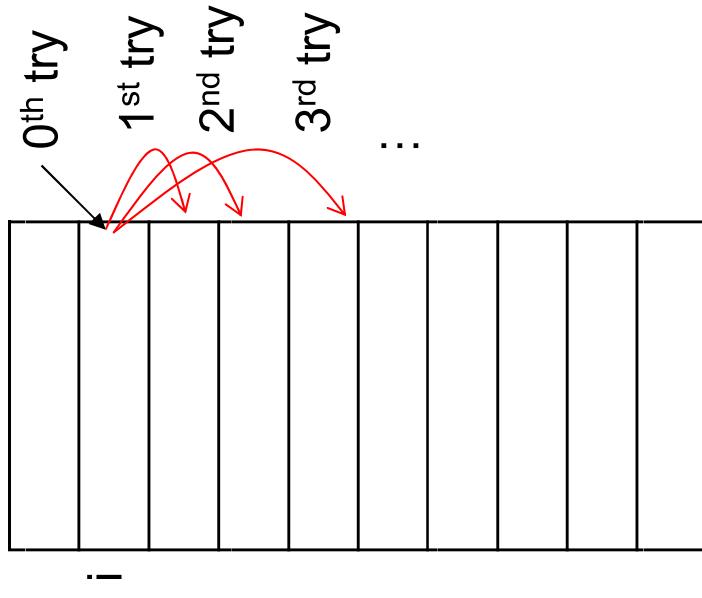
Double Hashing: Analysis

- Imperative that TableSize is prime
 - E.g., insert 23 into previous table
- Empirical tests show double hashing close to random hashing
- Extra hash function takes extra time to compute

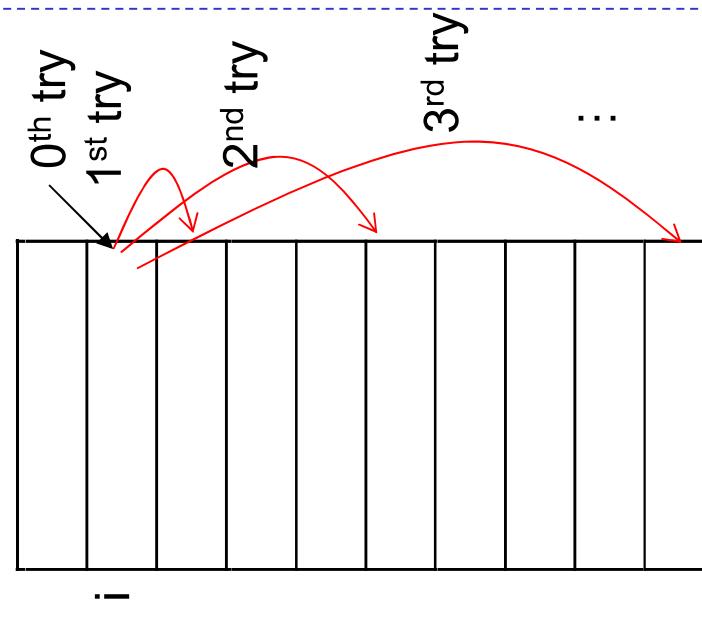
Probing Techniques - review



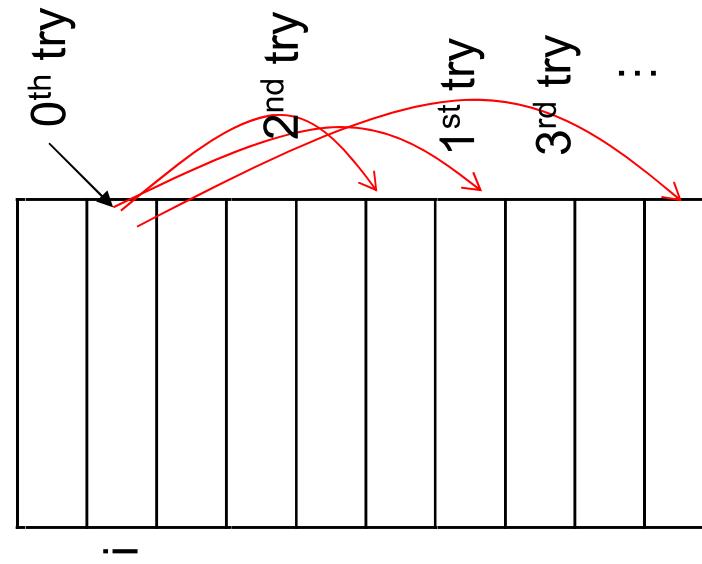
Linear probing:



Quadratic probing:



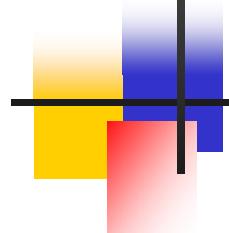
Double hashing*:



Rehashing

- Increases the size of the hash table when load factor becomes “too high” (defined by a cutoff)
 - Anticipating that prob(collisions) would become higher
- Typically expand the table to twice its size (but still prime)
- Need to reinsert all existing elements into new hash table

Rehashing Example



$$h(x) = x \bmod 7$$
$$\lambda = 0.57$$

0	6
1	15
2	
3	24
4	
5	
6	13

$$h(x) = x \bmod 17$$
$$\lambda = 0.29$$

Rehashing →
Insert 23 ↓

0	6
1	15
2	23
3	24
4	
5	
6	13

$$\lambda = 0.71$$

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	13
14	
15	15
16	
17	

Rehashing Analysis

- Rehashing takes time to do N insertions
- Therefore should do it infrequently
 - Specifically
 - Must have been $N/2$ insertions since last rehash
- Amortizing the $O(N)$ cost over the $N/2$ prior insertions yields only constant additional time per insertion

Rehashing Implementation

- When to rehash
 - When load factor reaches some threshold
(e.g., $\lambda \geq 0.5$), OR
 - When an insertion fails
- Applies across collision handling schemes

Rehashing for Chaining

```
20      /**
21       * Rehashing for separate chaining hash table.
22     */
23     void rehash( )
24     {
25       vector<list<HashedObj>> oldLists = theLists;
26
27       // Create new double-sized, empty table
28       theLists.resize( nextPrime( 2 * theLists.size( ) ) );
29       for( int j = 0; j < theLists.size( ); j++ )
30         theLists[ j ].clear( );
31
32       // Copy table over
33       currentSize = 0;
34       for( int i = 0; i < oldLists.size( ); i++ )
35     {
36         list<HashedObj>::iterator itr = oldLists[ i ].begin( );
37         while( itr != oldLists[ i ].end( ) )
38           insert( *itr++ );
39     }
40   }
```

Rehashing for Quadratic Probing

```
1  /**
2   * Rehashing for quadratic probing hash table.
3   */
4  void rehash( )
5  {
6      vector<HashEntry> oldArray = array;
7
8      // Create new double-sized, empty table
9      array.resize( nextPrime( 2 * oldArray.size( ) ) );
10     for( int j = 0; j < array.size( ); j++ )
11         array[ j ].info = EMPTY;
12
13     // Copy table over
14     currentSize = 0;
15     for( int i = 0; i < oldArray.size( ); i++ )
16         if( oldArray[ i ].info == ACTIVE )
17             insert( oldArray[ i ].element );
18 }
```

Hash Tables in C++ STL

- Hash tables not part of the C++ Standard Library
- Some implementations of STL have hash tables (e.g., SGI's STL)
 - **hash_set**
 - **hash_map**

Hash Set in STL

```
#include <hash_set>

struct eqstr
{
    bool operator() (const char* s1, const char* s2) const
    {
        return strcmp(s1, s2) == 0;
    }
};

void lookup(const hash_set<const char*, hash<const char*>, eqstr>& set,
            const char* word)
{
    hash_set<const char*, hash<const char*>, eqstr>::const_iterator it
        = set.find(word);
    cout << word << ":" <<
        << (it != set.end() ? "present" : "not present")
        << endl;
}

int main()
{
    hash_set<const char*, hash<const char*>, eqstr> set;
    set.insert("kiwi");
    lookup(set, "kiwi");
}
```

Key

Key equality test

Hash fn

Cpt S 223. School of EECS, WSU

Hash Map in STL

```
#include <hash_map>

struct eqstr
{
    bool operator() (const char* s1, const char* s2) const
    {
        return strcmp(s1, s2) == 0;
    }
};

int main()
{
    hash_map<const char*, int, hash<const char*>, eqstr> months;
    months["january"] = 31;
    months["february"] = 28;
    ...
    months["december"] = 31;
    cout << "january -> " << months["january"] << endl;
}
```

Internally
treated
like insert
(or overwrite
if key
already present)

Key equality test

Hash fn

Data

Problem with Large Tables

- What if hash table is too large to store in main memory?
- Solution: Store hash table on disk
 - Minimize disk accesses
 - But...
 - Collisions require disk accesses
 - Rehashing requires a lot of disk accesses

Solution: Extendible Hashing

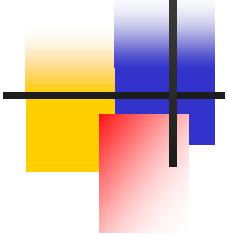
Hash Table Applications

- Symbol table in compilers
- Accessing tree or graph nodes by name
 - E.g., city names in Google maps
- Maintaining a transposition table in games
 - Remember previous game situations and the move taken (avoid re-computation)
- Dictionary lookups
 - Spelling checkers
 - Natural language understanding (word sense)
 - Heavily used in text processing languages
 - E.g., Perl, Python, etc.

Summary

- Hash tables support fast insert and search
 - $O(1)$ average case performance
 - Deletion possible, but degrades performance
- Not suited if ordering of elements is important
- Many applications

Points to remember - Hash tables



- Table size prime
- Table size much larger than number of inputs
(to maintain λ closer to 0 or < 0.5)
- Tradeoffs between chaining vs. probing
- Collision chances decrease in this order:
linear probing => quadratic probing =>
{random probing, double hashing}
- Rehashing required to resize hash table at a
time when λ exceeds 0.5
- Good for searching. Not good if there is some
order implied by ~~hash~~ table